

1. Overview

The Rebellions RBLN-CA12 accelerator card is a full height, full-length (FHFL), single-slot 10.5-inch PCI Express Gen5 equipped with ATOM™, Rebellions' System-on-Chip (SoC) designed for deep learning inference. The card is passively cooled.

RBLN-CA12 supports AI inference tasks for various neural network models, including Convolutional Neural Network, Recurrent Neural Network, and Transformers. The card mainly targets inference acceleration on edge-computing platforms as well as High Performance Computing (HPC) workloads. It can be deployed in cloud data center for scalable AI applications.